

Lecture 2: Samples, standard errors and confidence intervals

Empirical Methods for Law and Economics | EDLE

Marco Rosso
University of Bologna

Friday 9 October 2026

Today's question

If we had observed a different set of cases,
how different would our estimate be?

- By the end of today you will be able to:
 - distinguish a population parameter from a sample statistic;
 - explain the difference between a standard deviation and a standard error;
 - interpret a 95% confidence interval correctly, and state the assumptions behind it;
 - recognise when observations are not independent, as in “A Fine Is a Price”.

Retrieval: three minutes, on your own

1. Why is a court automation score of 0 different from a missing value?
2. The mean enforcement time was 650 days, the median 575. What does that tell you?
3. What did the SD of 304 days describe?

From describing to estimating


- Thursday: we described *all* 191 economies in the file. Nothing was estimated.
- Typical legal research: a court registry holds 10,000 closed commercial cases; we can read and code only 25 or 100 files.
- We want the average duration of *all* 10,000 cases. We only see a sample.
- Two different problems:
 - **sampling uncertainty**: another random sample would give a different mean;
 - **selection**: the files we can read may differ systematically from the others (e.g. only cases that reached judgment).
- Statistics quantifies the first. A larger sample does not cure the second.

Parameter, statistic, estimator

	Population (fixed, unknown)	Sample (varies)
Mean	μ	$\bar{y}$
Standard deviation	σ	s
Proportion	p	$\hat{p}$

- An **estimator** is a rule (“average the sampled durations”); an **estimate** is the number it gives in one sample.
- Same rule, different sample $\Rightarrow$ different estimate.

An artificial population where we know the truth



- Simulated case durations: right-skewed like Thursday's data; mean $\mu = 300$ days, SD $\sigma = 150$ days.
- Because we built it, we know μ . In real research we never do.

Standard deviation versus standard error

- **SD** (σ , estimated by s): how much individual cases differ from one another.
- **SE**: how much the *sample mean* would differ across repeated samples.

$$SE(\bar{y}) = \frac{\sigma}{\sqrt{n}}, \quad \text{estimated by } \frac{s}{\sqrt{n}}.$$

- With $\sigma = 150$: $n = 25 \Rightarrow 150/5 = 30$ days; $n = 100 \Rightarrow 150/10 = 15$ days.
- Quadrupling n halves the SE. The cases themselves do not become more similar: σ stays 150.

Why $\sqrt{n}$? (and why independence matters)

- If $y_1, \dots, y_n$ are independent with variance σ^2 :

$$\text{Var}\left(\sum_i y_i\right) = n\sigma^2 \quad \Rightarrow \quad \text{Var}(\bar{y}) = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n} \quad \Rightarrow \quad SE = \frac{\sigma}{\sqrt{n}}.$$

- The first step uses independence: positive and negative deviations of different cases partly cancel.
- If observations move together (same court, same judge, same centre), they cancel less, and $\sigma/\sqrt{n}$ is too optimistic.


Guided Stata: the supplied simulation

```
do "do/00_setup.do"  
do "do/ready_to_run/02_sampling_demo.do"  
use "results/tables/02_sampling_results.dta", clear  
tabstat samplemean sd se covers_truth, by(n) statistics(mean sd)
```

- The program draws 2,000 samples each of size $n = 5, 25$ and 100 and records, for each sample, its mean, SD, SE and 95% interval.
- One row of the results file = one simulated **sample**, not one case.
- You do not need to write simulation code. You need to explain the output.

What the simulation shows

2,000 sample means for each sample size (simulated)



	$n = 5$	$n = 25$	$n = 100$
Theoretical SE, $150/\sqrt{n}$	67.1	30.0	15.0
SD of the 2,000 sample means (approx.)	67	30	15

The central limit theorem, informally

- For independent draws with finite variance, the distribution of $\bar{y}$ becomes approximately normal as n grows, *whatever* the shape of individual outcomes.
- That is why we can build intervals with normal or t critical values.
- “Approximately” and “as n grows” matter: with $n = 5$ from a skewed population the approximation is still poor.
- With dependent observations the mean can still be approximately normal, but its SE is no longer $\sigma/\sqrt{n}$. Extremely heavy tails are a further problem.

A 95% confidence interval

$$\bar{y} \pm t_{n-1, 0.975} \times \frac{s}{\sqrt{n}}$$

- Centre: the estimate. Width: the estimated SE times a critical value.
- The t value (above 1.96, close to it once $n \geq 25$; 2.776 for $n = 5$) allows for the fact that s is estimated.
- Example: $n = 25$, $\bar{y} = 312$, $s = 150$: $SE = 30$, $t_{24} = 2.064$,


$$312 \pm 2.064 \times 30 = [250.1, 373.9] \text{ days.}$$

- With $n = 100$ and the same $\bar{y}$ and s : $SE = 15$, $t_{99} = 1.984$, interval $[282.2, 341.8]$.

What “95%” means

- The true mean (300) is fixed. The *intervals* change from sample to sample.
- If we repeated the sampling many times, about 95% of the intervals built this way would contain the true mean.
- It is a property of the **procedure**. A single realised interval either contains 300 or not; in real research we cannot tell which.
- It does *not* say that 95% of cases fall inside the interval, nor that the true mean has a 95% probability of lying in *this* interval.

Forty intervals from forty samples



Simulated: $n = 25$ per sample, true mean 300. Your Stata run will show different intervals.

Coverage depends on assumptions

	$n = 5$	$n = 25$	$n = 100$
Share of intervals containing 300 (approx.)	0.92	0.935	0.945

- “95%” is guaranteed only if the assumptions hold (independent random sampling; for small n , normal outcomes).
- With skewed outcomes and small samples the actual coverage falls short; it improves as n grows.
- With 2,000 replications the simulation itself is noisy by about ± 1 point (95% band): a run of 0.94 instead of 0.945 is not evidence of anything.

What a confidence interval cannot repair

- Measurement: a precise estimate of a badly measured outcome.
- Selection: the sample does not represent the population of interest.
- Dependence: repeated measurements on the same units treated as independent.
- Substance: is a difference of 1 day or of 40 days legally or economically important? Report intervals in original units.

- Ten private day-care centres in Haifa, observed for 20 weeks. Outcome: number of parents arriving late to collect their child, per centre and week.
- At the start of week 5 a fine (NIS 10 per late child) is introduced in 6 centres *selected randomly*; 4 centres serve as comparison.
- At the start of week 17 the fine is removed, without explanation.
- **Table 1:** one row per centre (with the number of children enrolled); one cell = late arrivals in that centre in one week.
- **Figure 1:** for each week, the average count across the centres of each group.

Discussion (30 minutes)

5 minutes alone, 10 in pairs, 10 whole class, 5 synthesis.

1. What does one row of Table 1 record, and what does one cell count? What does Figure 1 average?
2. Which centres receive the fine, how were they chosen, and in which weeks is it in force?
3. Why are the 200 centre-week counts not 200 independent observations of the effect of the fine?
4. Describe the pattern in Figure 1 in two sentences, without explaining *why* it happens.

From the $\sqrt{n}$ formula to the paper

- The fine was assigned to **centres**: 6 versus 4. The same centre is measured week after week.
- Weeks within a centre share parents, staff and routines, so they move together.
- Applying $s/\sqrt{200}$ as if there were 200 independent observations would overstate precision. The number of independently assigned units is 10.
- Constructive referee comment: “Please clarify how repeated weekly observations within centres are accounted for in the reported tests.”

Laboratory in pairs (40 minutes)

Template: `do/exercises/02_sampling_and_inference.do`.

1. Run the supplied simulation; say what one row of the results file is.
2. Compare the average SD, the SD of the sample means and the average SE for $n = 5, 25, 100$.
3. Look at the three histograms: where is the skew still visible?
4. Compute the coverage for each n ; explain why it is not exactly 0.95.
5. Extension: predict the SD of the sample means for $n = 400$, then check it (commands in the worksheet).
6. Write two paragraphs: SD vs SE with your numbers; a correct 95% interpretation.

Exit check

1. SD or SE: which describes variation across individual cases?
2. If n is multiplied by four, what happens to the SE, and what to the SD?
3. In “95% confidence”, what varies across repeated samples and what stays fixed?
4. In the day-care study, how many units were independently assigned to fine or no fine?

Next week

- Lesson 3: relationships between variables, hypothesis tests for means and proportions, association versus causation.
- Assessment: agree the paper or seminar for your referee report with me **by meeting 3**.